
BST 210 Lab: Week 9

Multinomial and Ordinal Regression

Multinomial Regression: A Review

Logistic Regression

Let's start by taking another look at logistic regression. In logistic regression, we consider binary outcomes, where an individual either experienced an event (success, disease, death, etc.) or not. First, let's recall exactly how we specify our model:

$$\text{logit}(P(Y = 1|X)) = \beta_0 + \beta_1 X_1 + \dots \beta_p X_p$$

Equivalently,

$$P(Y = 1|X) = \frac{\exp(\beta_0 + \beta_1 X_1 + \dots \beta_p X_p)}{1 + \exp(\beta_0 + \beta_1 X_1 + \dots \beta_p X_p)} = \pi_1$$

Since $1 = P(Y = 1|X) + P(Y = 0|X)$, we also get that:

$$P(Y = 0|X) = 1 - \frac{\exp(\beta_0 + \beta_1 X_1 + \dots \beta_p X_p)}{1 + \exp(\beta_0 + \beta_1 X_1 + \dots \beta_p X_p)} = \frac{1}{1 + \exp(\beta_0 + \beta_1 X_1 + \dots \beta_p X_p)} = \pi_0$$

We fit a logistic regression model and estimate parameters by using maximum likelihood. The likelihood can be seen as how likely we were to observe the data from our sample, given our chosen probability model and parameter values. For a binary outcome, the likelihood looks like:

$$\begin{aligned} L(\beta_0, \beta_1, \dots, \beta_p) &= \prod_{i=1}^n \pi_{i1}^{I(Y_i=1)} \pi_{i0}^{I(Y_i=0)} \\ &= \prod_{i=1}^n \left(\frac{\exp(\beta_0 + \beta_1 X_{i1} + \dots \beta_p X_{ip})}{1 + \exp(\beta_0 + \beta_1 X_{i1} + \dots \beta_p X_{ip})} \right)^{y_i} \left(\frac{1}{1 + \exp(\beta_0 + \beta_1 X_{i1} + \dots \beta_p X_{ip})} \right)^{1-y_i} \end{aligned}$$

This is solely a function of our beta coefficients and can be maximized using iterative computation techniques as our software packages do.

Multinomial Regression

Now let's consider extending what we know about logistic regression using binary outcomes to multinomial regression. Here, instead of Y only taking on two possible values (0 or 1), Y can take on c different values (from 1 up to c). Because each individual must have only one outcome, we can say:

$$\sum_{j=1}^c \pi_j = 1$$

where $\pi_k = P(Y = k|X)$. In multinomial regression, our model looks like (use category 1 as the baseline):

$$\log\left(\frac{\pi_k}{\pi_1}\right) = \beta_{k0} + \beta_{k1}X_1 + \dots + \beta_{kp}X_p \quad (k \in (2, \dots, c))$$

Using our model and the fact that the π_k must sum to 1, we can show:

$$\pi_1 = \frac{1}{1 + \sum_{j=2}^c \exp(\beta_{j0} + \beta_{j1}X_1 + \dots + \beta_{jp}X_p)}$$

(See Appendix 1 for detailed derivations.)

Once again, using our model, for $k = 2, \dots, c$, we can write:

$$\begin{aligned} \pi_k &= \pi_1 \times \exp(\beta_{k0} + \beta_{k1}X_1 + \dots + \beta_{kp}X_p) \\ &= \frac{\exp(\beta_{k0} + \beta_{k1}X_1 + \dots + \beta_{kp}X_p)}{1 + \sum_{j=2}^c \exp(\beta_{j0} + \beta_{j1}X_1 + \dots + \beta_{jp}X_p)} \end{aligned}$$

Using these probabilities, we can now extend our likelihood from logistic regression to a multinomial likelihood:

$$L(\beta_0, \beta_1, \dots, \beta_p) = \prod_{i=1}^n \pi_{i1}^{I(Y_i=1)} \times \dots \times \pi_{ic}^{I(Y_i=c)}$$

Replacing the π_k s with their formula involving β s, we can simply use iterative methods once again to find our maximum likelihood estimates.

Multinomial Example

A study examined the relative efficacy of penicillin and spectinomycin in treating gonorrhea. Three treatments were considered: (1) penicillin, (2) spectinomycin [Low Dose], and (3) spectinomycin [High Dose]. Three possible responses were recorded: (1) positive smear, (2) negative smear, positive culture, and (3) negative smear, negative culture. The results of the study are presented in the figure below:

Treatment	Response			Total
	+ Smear	- Smear + Culture	- Smear - Culture	
Penicillin	40	30	130	200
Spectinomycin, low dose	10	20	70	100
Spectinomycin, high dose	15	40	45	100
Total	65	90	245	400

Figure 1: Results from Gonorrhea Study

Using the following Stata outputs, answer the following questions:

1. Write out the full model.

Multinomial Model Fit

```
. *****
. * BST 210 Week 9 Lab *
. *****
. * Input data
. clear all
. use "data_lab9.dta"

. *Multinomial logistic regression
. mlogit outcome i.treatment [fweight = count], baseoutcome(1)

Iteration 0:   log likelihood = -372.45951
Iteration 1:   log likelihood = -358.97348
Iteration 2:   log likelihood = -358.51621
Iteration 3:   log likelihood = -358.51603
Iteration 4:   log likelihood = -358.51603

Multinomial logistic regression      Number of obs      =          400
LR chi2(4)          =          27.89
Prob > chi2         =          0.0000
Log likelihood = -358.51603          Pseudo R2           =          0.0374

-----+-----
outcome   |      Coef.   Std. Err.      z    P>|z|     [95% Conf. Interval]
-----+-----
1         |      (base outcome)
-----+-----
2         |
treatment |
Spect-Low |   .9808293   .4564355     2.15   0.032   .0862322   1.875426
Spect-Hight |  1.268511   .3872983     3.28   0.001   .5094205   2.027602
|
_cons     |  -.2876821   .2415229    -1.19   0.234  -.7610583   .1856942
-----+-----
3         |
treatment |
Spect-Low |   .7672552   .3833771     2.00   0.045   .0158498   1.518661
Spect-Hight | -.0800427   .348685    -0.23   0.818  -.7634527   .6033672
|
_cons     |   1.178655   .1808101     6.52   0.000   .8242736   1.533036
-----+-----
```

```
* We can directly get RRRs instead of beta coefficients with the 'rrr' option
. mlogit outcome i.treatment [fweight = count], baseoutcome(1) rrr
```

```
Iteration 0:   log likelihood = -372.45951
Iteration 1:   log likelihood = -358.97348
Iteration 2:   log likelihood = -358.51621
Iteration 3:   log likelihood = -358.51603
Iteration 4:   log likelihood = -358.51603
```

```
Multinomial logistic regression      Number of obs      =          400
LR chi2(4)          =          27.89
Prob > chi2         =          0.0000
Log likelihood = -358.51603          Pseudo R2           =          0.0374
```

outcome	RRR	Std. Err.	z	P> z	[95% Conf. Interval]	
1	(base outcome)					
2						
treatment						
Spect-Low	2.666667	1.217161	2.15	0.032	1.090059	6.5236
Spect-Hight	3.555556	1.377061	3.28	0.001	1.664327	7.595851
_cons	.75	.1811422	-1.19	0.234	.4671717	1.204054
3						
treatment						
Spect-Low	2.153846	.8257354	2.00	0.045	1.015976	4.566105
Spect-Hight	.9230769	.321863	-0.23	0.818	.4660545	1.828265
_cons	3.25	.587633	6.52	0.000	2.280224	4.632221

- 5

4. Using the model output, calculate the fitted probabilities of the outcome for a population having the Penicillin treatment ($P(Y = 1|Pen)$, $P(Y = 2|Pen)$, $P(Y = 3|Pen)$).

Fitted Values

```
. predict fitm1 fitm2 fitm3
(option pr assumed; predicted probabilities)

. list outcome treatment fitm1 fitm2 fitm3
```

	outcome	treatment	fitm1	fitm2	fitm3
1.	1	Penicillin	.2	.15	.65
2.	1	Spect-Low	.1	.2	.7
3.	1	Spect-Hight	.15	.4	.45
4.	2	Penicillin	.2	.15	.65
5.	2	Spect-Low	.1	.2	.7
6.	2	Spect-Hight	.15	.4	.45
7.	3	Penicillin	.2	.15	.65
8.	3	Spect-Low	.1	.2	.7
9.	3	Spect-Hight	.15	.4	.45

5. How do the fitted probabilities from the model compare to the observed proportions in the data table? Why is this the case?

Some More Stata Results:

```
. *If we don't specify the reference group, Stata automatically uses
. *the most frequent outcome as the reference group.

. mlogit outcome i.treatment [fweight = count]
```

Iteration 0: log likelihood = -372.45951
Iteration 1: log likelihood = -358.97348
Iteration 2: log likelihood = -358.51621
Iteration 3: log likelihood = -358.51603
Iteration 4: log likelihood = -358.51603

Multinomial logistic regression Number of obs = 400
LR chi2(4) = 27.89
Prob > chi2 = 0.0000
Log likelihood = -358.51603 Pseudo R2 = 0.0374

outcome	Coef.	Std. Err.	z	P> z	[95% Conf. Interval]	
1						
treatment						
Spect-Low	-.7672552	.3833771	-2.00	0.045	-1.518661	-.0158498
Spect-High	.0800427	.348685	0.23	0.818	-.6033672	.7634527
_cons	-1.178655	.1808101	-6.52	0.000	-1.533036	-.8242736
2						
treatment						
Spect-Low	.2135741	.3245171	0.66	0.510	-.4224677	.849616
Spect-High	1.348554	.2970654	4.54	0.000	.7663165	1.930792
_cons	-1.466337	.2025479	-7.24	0.000	-1.863324	-1.069351
3	(base outcome)					

```
. estat ic
```

Akaike's information criterion and Bayesian information criterion

Model	Obs	ll(null)	ll(model)	df	AIC	BIC
.	400	-372.4595	-358.516	6	729.0321	752.9808

Note: N=Obs used in calculating BIC; see [R] BIC note.

Ordinal Logistic Regression Review

For ordinal logistic regression, we are going to make the proportional odds assumption, giving us the following model specification:

$$\log \left(\frac{P(Y \geq j)}{P(Y < j)} \right) = \alpha_j + \beta_1 X_1 + \dots + \beta_p X_p$$

With the proportional odds assumption, we assume that for each cutpoint j , we are only shifting our logistic curve by a constant, namely, the α_j s. This in turn implies that the effect of a certain covariate β_1 is constant across different cutpoints in our outcome.

NOTE: Different packages will specify this model in different ways. For example, many will instead have:

$$\log \left(\frac{P(Y \leq j)}{P(Y > j)} \right) = \alpha_j + \beta_1 X_1 + \dots + \beta_p X_p$$

Also, some packages will specify the model using negative intercepts. Please pay attention to how your package is specifying the model when it comes to interpretation.

Ordinal Logistic Regression Example

We'll be using the same data as the previous example, but instead, take into account the ordinal nature of our outcome [Outcome Smear-positive is worse than Outcome Smear-negative/Culture-positive is worse than outcome Smear-negative/Culture-negative].

Using the following Stata output, answer the questions below:

1. Write out the proportional odds ordinal logistic model. How is this different than the multinomial model? Why would you use one over the other?

Proportional Odds

```
. *Ordinal logistic regression
. ologit outcome i.treatment [fweight = count]
```

Iteration 0: log likelihood = -372.45951
Iteration 1: log likelihood = -367.41299
Iteration 2: log likelihood = -367.38955
Iteration 3: log likelihood = -367.38955

Ordered logistic regression Number of obs = 400
LR chi2(2) = 10.14
Prob > chi2 = 0.0063
Log likelihood = -367.38955 Pseudo R2 = 0.0136

outcome	Coef.	Std. Err.	z	P> z	[95% Conf. Interval]	
treatment						
Spect-Low	.3410902	.260078	1.31	0.190	-.1686534	.8508338
Spect-Hight	-.5194078	.2325509	-2.23	0.026	-.9751992	-.0636163
/cut1	-1.73064	.175232			-2.074088	-1.387192
/cut2	-.5251288	.1487098			-.8165947	-.2336628

```
* We can directly get ORs intead of beta coefficients with the 'or' option
. ologit outcome i.treatment [fweight = count], or
```

Iteration 0: log likelihood = -372.45951
Iteration 1: log likelihood = -367.41299
Iteration 2: log likelihood = -367.38955
Iteration 3: log likelihood = -367.38955

Ordered logistic regression Number of obs = 400
LR chi2(2) = 10.14
Prob > chi2 = 0.0063
Log likelihood = -367.38955 Pseudo R2 = 0.0136

outcome	Odds Ratio	Std. Err.	z	P> z	[95% Conf. Interval]	
treatment						
Spect-Low	1.40648	.3657946	1.31	0.190	.8448017	2.341598
Spect-Hight	.5948727	.1383382	-2.23	0.026	.3771172	.9383649
/cut1	-1.73064	.175232			-2.074088	-1.387192
/cut2	-.5251288	.1487098			-.8165947	-.2336628

```
. estat ic
```

Akaike's information criterion and Bayesian information criterion

Model	Obs	ll(null)	ll(model)	df	AIC	BIC
.	400	-372.4595	-367.3895	4	742.7791	758.745

Note: N=Obs used in calculating BIC; see [R] BIC note.

2. Provide brief interpretations β_1 .

Fitted Values

```
. predict fito1 fito2 fito3
(option pr assumed; predicted probabilities)

. list outcome treatment fito1 fito2 fito3
```

	outcome	treatment	fito1	fito2	fito3
1.	1	Penicillin	.1505057	.221148	.6283463
2.	1	Spect-Low	.111875	.1841667	.7039583
3.	1	Spect-Hight	.2294831	.2690867	.5014302
4.	2	Penicillin	.1505057	.221148	.6283463
5.	2	Spect-Low	.111875	.1841667	.7039583
6.	2	Spect-Hight	.2294831	.2690867	.5014302
7.	3	Penicillin	.1505057	.221148	.6283463
8.	3	Spect-Low	.111875	.1841667	.7039583
9.	3	Spect-Hight	.2294831	.2690867	.5014302

3. How do the fitted probabilities from the ordinal model compare to the observed proportions in the data? Why?

4. How could you assess if the proportional odds assumption is violated? What are your conclusions in this case?

Relaxing the Proportional Odds Assumption

```
. gologit2 outcome i.treatment [fweight = count], rrr
```

Generalized Ordered Logit Estimates		Number of obs	=	400
LR chi2(4)	=	27.89		
Prob > chi2	=	0.0000		
Log likelihood = -358.51603		Pseudo R2	=	0.0374

outcome	RRR	Std. Err.	z	P> z	[95% Conf. Interval]	
1						
treatment						
Spect-Low	2.250001	.8489427	2.15	0.032	1.074033	4.713544
Spect-Hight	1.416667	.4691744	1.05	0.293	.7402279	2.711254
_cons	3.999999	.7071064	7.84	0.000	2.828702	5.656299
2						
treatment						
Spect-Low	1.25641	.3314568	0.87	0.387	.7491586	2.107119
Spect-Hight	.4405594	.1100359	-3.28	0.001	.2700256	.7187932
_cons	1.857143	.2753212	4.18	0.000	1.388848	2.483338


```
. estat ic
```

Akaike's information criterion and Bayesian information criterion

Model	Obs	ll(null)	ll(model)	df	AIC	BIC
.	400	-372.4595	-358.516	6	729.0321	752.9808

Note: N=Obs used in calculating BIC; see [R] BIC note.

Testing Proportional Odds Assumption

```

. gen spectlow = (treatment == 2)

. gen specthigh = (treatment == 3)

. omodel logit outcome spectlow specthigh [fweight = count]

Iteration 0:   log likelihood = -372.45951
Iteration 1:   log likelihood = -367.41299
Iteration 2:   log likelihood = -367.38955
Iteration 3:   log likelihood = -367.38955

Ordered logit estimates               Number of obs   =           400
LR chi2(2)      =           10.14
Prob > chi2     =           0.0063
Log likelihood = -367.38955           Pseudo R2       =           0.0136
-----
outcome        |      Coef.   Std. Err.      z    P>|z|     [95% Conf. Interval]
-----+-----
spectlow       |   .3410902   .260078     1.31   0.190    -.1686534   .8508338
specthigh      |  -.5194078   .2325509    -2.23   0.026    -.9751992  -.0636163
-----+-----
_cut1          |  -1.73064    .175232                (Ancillary parameters)
_cut2          |  -.5251288   .1487098
-----+-----

Approximate likelihood-ratio test of proportionality of odds
across response categories:
chi2(2) =      18.16
Prob > chi2 =     0.0001

. *use "ssc install omodel" if you don't have this package

```

Food for thought: can you think of other ways to assess the assumption (not necessarily a formal statistical test)? See Appendix 3 for one possible way.

5. Which model should we use for this data? Why?

Appendix

1. Derivation of π_1

Use our multinomial logistic regression model and the fact that the π_k must sum to 1, we can show:

$$\pi_1 = \frac{1}{1 + \sum_{j=2}^c \exp(\beta_{j0} + \beta_{j1}X_1 + \dots + \beta_{jp}X_p)}$$

Let's denote $\mathbf{X}\beta_k = \beta_{k0} + \beta_{k1}X_1 + \dots + \beta_{kp}X_p$ (note this is actually the matrix form), and our model becomes:

$$\begin{aligned} \log\left(\frac{\pi_k}{\pi_1}\right) &= \mathbf{X}\beta_k \\ \Leftrightarrow \frac{\pi_k}{\pi_1} &= \exp(\mathbf{X}\beta_k) \\ \Leftrightarrow \pi_k &= \pi_1 \exp(\mathbf{X}\beta_k) \end{aligned}$$

Now let's return our attention to the sum of the π s:

$$\begin{aligned} 1 &= \pi_1 + \pi_2 + \dots + \pi_c \\ \Leftrightarrow 1 &= \pi_1 + \pi_1 \exp(\mathbf{X}\beta_1) + \dots + \pi_1 \exp(\mathbf{X}\beta_c) \\ \Leftrightarrow 1 &= \pi_1 \left(1 + \sum_{j=2}^c \exp(\mathbf{X}\beta_j) \right) \\ \Leftrightarrow \pi_1 &= \frac{1}{1 + \sum_{j=2}^c \exp(\mathbf{X}\beta_j)} \end{aligned}$$

2. Hand calculation of RRR

Use the data table, compute (by hand), the relative risk ratios for having outcome 2 versus outcome 1 comparing people given Specitinomycin Low Dose to people given Penicillin, and compare them to the results found in the package output.

$$RRR = \frac{\frac{P(Y=2|LD)}{P(Y=1|LD)}}{\frac{P(Y=2|Pen)}{P(Y=1|Pen)}} = \frac{\frac{\frac{20}{100}}{\frac{100}{30}}}{\frac{\frac{200}{40}}{\frac{200}{200}}} = \frac{\frac{20}{30}}{\frac{10}{40}} = \frac{8}{3}$$

This is the same as what we got from the package.

3. Assess proportional odds assumption

We can compare two logistic regression models as follows:

```
. gen r1 = (outcome > 1)
```

```
. gen r2 = (outcome > 2)
```

```
. logistic r1 i.treatment [fweight = count]
```

```
Logistic regression               Number of obs   =           400
LR chi2(2)                       =           5.31
Prob > chi2                      =           0.0701
Log likelihood = -174.85969       Pseudo R2      =           0.0150
```

```
-----
```

r1	Odds Ratio	Std. Err.	z	P> z	[95% Conf. Interval]	
treatment						
Spect-Low		2.25	.8489424	2.15	0.032	1.074033 4.713542
Spect-Hight		1.416667	.4691742	1.05	0.293	.7402275 2.711253
_cons		4	.7071068	7.84	0.000	2.828703 5.656302

```
-----
```

```
. logistic r2 i.treatment [fweight = count]
```

```
Logistic regression               Number of obs   =           400
LR chi2(2)                       =          15.31
Prob > chi2                      =           0.0005
Log likelihood = -259.38964       Pseudo R2      =           0.0287
```

```
-----
```

r2	Odds Ratio	Std. Err.	z	P> z	[95% Conf. Interval]	
treatment						
Spect-Low		1.25641	.3314568	0.87	0.387	.7491587 2.107119
Spect-Hight		.4405594	.1100359	-3.28	0.001	.2700256 .7187933
_cons		1.857143	.2753212	4.18	0.000	1.388848 2.483338

```
-----
```

What do we notice about the coefficients in these two models? Do they tell you anything?